

```
!pip install torch torchvision opencv-python tqdm scikit-image matplotlib
```

```
!pip install -U scikit-image --quiet
```

```
!pip install --upgrade scikit-image
```

```
Requirement already satisfied: torch in /usr/local/lib/python3.11/dist-packages (2.6.0+cu124)
Requirement already satisfied: torchvision in /usr/local/lib/python3.11/dist-packages (0.21.0+cu124)
Requirement already satisfied: opencv-python in /usr/local/lib/python3.11/dist-packages (4.11.0.86)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (4.67.1)
Requirement already satisfied: scikit-image in /usr/local/lib/python3.11/dist-packages (0.25.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (3.10.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from torch) (3.18.0)
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.11/dist-packages (from torch) (4.14.0)
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch) (3.5)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.11/dist-packages (from torch) (3.1.6)
Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from torch) (2025.3.2)
Collecting nvidia-cuda-nvrtc-cu12==12.4.127 (from torch)
  Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-runtime-cu12==12.4.127 (from torch)
  Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cuda-cupti-cu12==12.4.127 (from torch)
  Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cudnn-cu12==9.1.0.70 (from torch)
  Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cublas-cu12==12.4.5.8 (from torch)
  Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cufft-cu12==11.2.1.3 (from torch)
  Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-curand-cu12==10.3.5.147 (from torch)
  Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Collecting nvidia-cusolver-cu12==11.6.1.9 (from torch)
  Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Collecting nvidia-cusparselt-cu12==0.6.2 (from torch)
  Downloading nvidia_cusparselt_cu12-0.6.2-py3-none-manylinux2014_x86_64.whl.metadata (1.6 kB)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /usr/local/lib/python3.11/dist-packages (from torch) (2.21.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch) (12.4.127)
Collecting nvidia-nvjitlink-cu12==12.4.127 (from torch)
  Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl.metadata (1.5 kB)
Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/dist-packages (from torch) (3.2.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (from torch) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch) (1.3.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from torchvision) (2.0.2)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.11/dist-packages (from torchvision) (11.2.1)
Requirement already satisfied: scipy>=1.11.4 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (1.15.3)
Requirement already satisfied: imageio<2.35.0,>=2.33 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (2.37.0)
Requirement already satisfied: tifffile>=2022.8.12 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (2025.6.11)
Requirement already satisfied: packaging>=21 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (24.2)
Requirement already satisfied: lazy-loader>=0.4 in /usr/local/lib/python3.11/dist-packages (from scikit-image) (0.4)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3.2)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (4.58.4)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.4.8)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2.3)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (2.9.0.post0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7->matplotlib) (1.17.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages (from Jinja2->torch) (3.0.2)
Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl (363.4 MB)
 363.4/363.4 MB 3.9 MB/s eta 0:00:00
Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (13.8 MB)
 13.8/13.8 MB 67.0 MB/s eta 0:00:00
Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (24.6 MB)
 24.6/24.6 MB 36.7 MB/s eta 0:00:00
```

```
# PHASE 1: Setup & Dataset Download
```

```
!pip install opencv-python scikit-image einops --quiet
```

```
import os, zipfile, glob, cv2, requests
```

```
import numpy as np
```

```
from tqdm import tqdm
```

```
from matplotlib import pyplot as plt
```

```
from skimage.metrics import structural_similarity as ssim
```

```
from PIL import Image
```

```
# Paths
```

```
BASE_DIR = "/content/sharpeness_data"
```

```
HR_DIR = os.path.join(BASE_DIR, "HR")
```

```
LR_DIR = os.path.join(BASE_DIR, "LR")
```

```
os.makedirs(HR_DIR, exist_ok=True)
```

```
os.makedirs(LR_DIR, exist_ok=True)
```

```
# 100 images from DIV2K (HR)
DIV2K_URL = "http://data.vision.ee.ethz.ch/cv1/DIV2K/DIV2K_train_HR.zip"
zip_path = "/content/DIV2K_train_HR.zip"

if not os.path.exists(zip_path):
    print("Downloading DIV2K subset...")
    r = requests.get(DIV2K_URL, stream=True)
    with open(zip_path, 'wb') as f:
        for chunk in r.iter_content(chunk_size=1024):
            if chunk: f.write(chunk)
    print("Extracting...")
    with zipfile.ZipFile(zip_path, 'r') as zip_ref:
        zip_ref.extractall("/content")

# Preprocess images: Simulate blur
for img_path in tqdm(sorted(glob.glob("/content/DIV2K_train_HR/*.png"))[:100]):
    img = cv2.imread(img_path)
    img = cv2.resize(img, (256, 256)) # normalize
    degraded = cv2.GaussianBlur(img, (5,5), 1) # blur

    fname = os.path.basename(img_path)
    cv2.imwrite(os.path.join(HR_DIR, fname), img)
    cv2.imwrite(os.path.join(LR_DIR, fname), degraded)

print("✅ Data ready: 100 LR-HR image pairs created.")
```

```
↻ Downloading DIV2K subset...
Extracting...
100%|██████████| 100/100 [00:10<00:00, 9.99it/s] ✅ Data ready: 100 LR-HR image pairs created.
```

```
import torch
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms

class SharpnessDataset(Dataset):
    def __init__(self, lr_dir, hr_dir):
        self.lr_paths = sorted(glob.glob(os.path.join(lr_dir, "*.png")))
        self.hr_paths = sorted(glob.glob(os.path.join(hr_dir, "*.png")))
        self.tf = transforms.ToTensor()

    def __len__(self):
        return len(self.lr_paths)

    def __getitem__(self, idx):
        lr = self.tf(Image.open(self.lr_paths[idx]).convert('RGB'))
        hr = self.tf(Image.open(self.hr_paths[idx]).convert('RGB'))
        return lr, hr

dataset = SharpnessDataset(LR_DIR, HR_DIR)
dataloader = DataLoader(dataset, batch_size=4, shuffle=True)

import torch.nn as nn
import torch.nn.functional as F

class ResidualBlock(nn.Module):
    def __init__(self, channels):
        super().__init__()
        self.block = nn.Sequential(
            nn.Conv2d(channels, channels, 3, padding=1),
            nn.ReLU(inplace=True),
            nn.Conv2d(channels, channels, 3, padding=1)
        )

    def forward(self, x):
        return x + self.block(x)

class SharpResNet(nn.Module):
    def __init__(self, blocks=5):
        super().__init__()
        self.entry = nn.Conv2d(3, 64, 3, padding=1)
        self.resblocks = nn.Sequential(*[ResidualBlock(64) for _ in range(blocks)])
        self.exit = nn.Conv2d(64, 3, 3, padding=1)
```

```

def forward(self, x):
    x = self.entry(x)
    x = self.resblocks(x)
    x = self.exit(x)
    return x

from torchvision.models import vgg16, VGG16_Weights

# VGG-based perceptual loss
vgg = vgg16(weights=VGG16_Weights.IMAGENET1K_V1).features[:16].eval().cuda()
for p in vgg.parameters():
    p.requires_grad = False

def perceptual_loss(pred, target):
    return F.mse_loss(vgg(pred), vgg(target))

# Custom SSIM loss
def gaussian(window_size, sigma):
    gauss = torch.Tensor([np.exp(-(x - window_size//2)**2 / float(2*sigma**2)) for x in range(window_size)])
    return gauss / gauss.sum()

def create_window(window_size, channel):
    _1D = gaussian(window_size, 1.5).unsqueeze(1)
    _2D = (_1D @ _1D.t()).float().unsqueeze(0).unsqueeze(0)
    return _2D.expand(channel, 1, window_size, window_size).contiguous()

def ssim(img1, img2, window_size=11):
    (_, c, h, w) = img1.size()
    window = create_window(window_size, c).to(img1.device)
    mu1 = F.conv2d(img1, window, padding=window_size//2, groups=c)
    mu2 = F.conv2d(img2, window, padding=window_size//2, groups=c)

    mu1_sq = mu1 ** 2
    mu2_sq = mu2 ** 2
    mu1_mu2 = mu1 * mu2


    sigma1_sq = F.conv2d(img1*img1, window, padding=window_size//2, groups=c) - mu1_sq
    sigma2_sq = F.conv2d(img2*img2, window, padding=window_size//2, groups=c) - mu2_sq
    sigma12 = F.conv2d(img1*img2, window, padding=window_size//2, groups=c) - mu1_mu2

    C1 = 0.01 ** 2
    C2 = 0.03 ** 2

    ssim_map = ((2 * mu1_mu2 + C1)*(2 * sigma12 + C2)) / ((mu1_sq + mu2_sq + C1)*(sigma1_sq + sigma2_sq + C2))
    return ssim_map.mean()

# ☒ Combined loss
def total_loss(pred, target):
    return (
        0.5 * F.mse_loss(pred, target) +
        0.3 * perceptual_loss(pred, target) +
        0.2 * (1 - ssim(pred, target))
    )

```

 Downloading: "<https://download.pytorch.org/models/vgg16-397923af.pth>" to /root/.cache/torch/hub/checkpoints/vgg16-397923af.pth
100%|██████████| 528M/528M [00:05<00:00, 92.3MB/s]

```

model = SharpResNet().cuda()
optimizer = torch.optim.Adam(model.parameters(), lr=1e-4)

```

```

for epoch in range(10):
    model.train()
    epoch_loss = 0
    for lr, hr in dataloader:
        lr, hr = lr.cuda(), hr.cuda()
        preds = model(lr)


        loss = total_loss(preds, hr)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        epoch_loss += loss.item()

```

```
print(f"Epoch {epoch+1}/10 - Loss: {epoch_loss / len(dataloader):.4f}")
```

```
Epoch 1/10 - Loss: 0.9506  
Epoch 2/10 - Loss: 0.5992  
Epoch 3/10 - Loss: 0.4654  
Epoch 4/10 - Loss: 0.3857  
Epoch 5/10 - Loss: 0.3270  
Epoch 6/10 - Loss: 0.2760  
Epoch 7/10 - Loss: 0.2234  
Epoch 8/10 - Loss: 0.1870  
Epoch 9/10 - Loss: 0.1726  
Epoch 10/10 - Loss: 0.1604
```

```
from skimage.metrics import structural_similarity as ssim
```

```
model.eval()  
with torch.no_grad():  
    for i in range(3):  
        lr, hr = dataset[i]  
        inp = lr.unsqueeze(0).cuda()  
        out = model(inp).clamp(0, 1).squeeze().cpu().permute(1, 2, 0).numpy()  
        orig = hr.permute(1, 2, 0).numpy()  
        blurred = lr.permute(1, 2, 0).numpy()  
  
        #  SSIM Score between output and ground truth  
        ssim_score = ssim(out, orig, data_range=1.0, channel_axis=-1)  
        print(f"SSIM Score for sample {i+1}: {ssim_score:.4f}")  
  
        plt.figure(figsize=(12,4))  
        plt.subplot(1,3,1); plt.imshow(blurred); plt.title("Blurred Input"); plt.axis("off")  
        plt.subplot(1,3,2); plt.imshow(out); plt.title("Enhanced Output"); plt.axis("off")  
        plt.subplot(1,3,3); plt.imshow(orig); plt.title("Ground Truth"); plt.axis("off")  
        plt.show()
```

SSIM Score for sample 1: 0.8476

Blurred Input



Enhanced Output



Ground Truth



```
import os
import glob
import time
import torch
import cv2

def measure_fps(model, device, test_image_path, num_runs=50):
    model.eval()
    img = cv2.imread(test_image_path)
    img = cv2.resize(img, (256, 256))
    input_tensor = torch.tensor(img).permute(2, 0, 1).unsqueeze(0).float() / 255.0
    input_tensor = input_tensor.to(device)

    # Warm-up
    for _ in range(10):
        with torch.no_grad():
            _ = model(input_tensor)

    start = time.time()
    for _ in range(num_runs):
        with torch.no_grad():
            _ = model(input_tensor)
    end = time.time()

    fps = num_runs / (end - start)
    print(f"✅ Model Inference Speed: {fps:.2f} FPS")
    return fps

# Set device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Ensure model is on the correct device
model = model.to(device)

# Run this after training
test_img_path = glob.glob(os.path.join(LR_DIR, "*.png"))[0] # Use first LR image
_ = measure_fps(model, device, test_img_path)
```

✅ Model Inference Speed: 113.99 FPS

